

Implementing Redis-Based Background Jobs in Photon (Rust)

Building a **fire-and-forget** background job system in Rust (inspired by Sidekiq) involves using Redis as a distributed queue and separate worker processes to process jobs. In Photon – a coroutine-based Rust web framework – this means the web server enqueues jobs (serialized as JSON) into Redis without blocking the request/response cycle, and dedicated Rust worker programs consume and execute these jobs. This ensures that time-consuming tasks run asynchronously, keeping the web app responsive ¹. Below is a guide on the architecture and implementation details, including job enqueueing, status tracking, example code, and API design for job status.

Architecture: Redis Queue + Rust Worker Processes

Redis-backed Queue: A common pattern for background jobs is to use Redis as a durable, fast queue. The web app pushes a job message (e.g. JSON data) onto a Redis list or stream, and one or more **worker processes** poll that queue to retrieve jobs. This is the same model used by Sidekiq in Ruby, which serializes job info into JSON and pushes it to a Redis queue ². Worker processes (or threads) continuously poll the Redis list (e.g. using `BLPOP` for blocking pop) and, when a job is found, deserialize the JSON and execute the task.

Fire-and-forget Execution: Once Photon enqueues the job message, it does **not wait** for the job to complete. The HTTP request can return immediately (usually with a `202 Accepted` status), while the job is handled in the background. Offloading work like this keeps the main application thread free to handle other requests ³. For example, instead of processing an expensive operation (sending emails, generating reports, etc.) during an HTTP request, Photon would enqueue a job to Redis and respond immediately, allowing a worker to handle the heavy lifting. This design improves responsiveness and throughput of the web server.

Multiple Workers & Scalability: Because jobs are queued in Redis, you can run multiple worker processes (potentially on different machines or cores) to consume the queue in parallel. This enables horizontal scaling – as load increases, you just add more workers reading from the same Redis queue ⁴. Each worker will pull jobs and process them independently, which is analogous to Sidekiq using multiple threads or processes to handle jobs concurrently ⁴. The queue ensures work is distributed and that each job is only handled once.

Existing Libraries vs Custom Implementation: There are Rust libraries that already implement this queue/worker pattern. For example, **Apalis** is a crate that provides background job handling with Redis support ⁵, and **rusty_sidekiq/sidekiq-rs** is a Sidekiq-compatible Rust implementation. These libraries let you define job types and workers, handling retries and more. However, it's also straightforward to implement a simple Redis-backed job queue yourself with `redis` and `serde` if you prefer a minimal custom solution. The core architecture remains: **Photon enqueues jobs to Redis, and separate Rust workers dequeue and execute them** ⁶.

Enqueuing Jobs in Photon (Non-Blocking Fire-and-Forget)

In Photon, when a request triggers a background task, the handler should **enqueue the job to Redis and return immediately**. Photon is coroutine-based (using the `may` library for stackful coroutines), so it can perform the Redis operation without blocking other requests. The coroutine can yield while waiting on Redis I/O, similar to `async/await`, meaning the server thread isn't stalled.

1. Define the Job Data: First, define a *job payload structure* and ensure it's serializable to JSON. For example, if we want to send welcome emails in the background, we might define:

```
use serde::{Serialize, Deserialize};
#[derive(Serialize, Deserialize)]
struct EmailJob {
    id: String,           // unique job ID
    to: String,
    subject: String,
    body: String
}
```

Here, `id` will be a unique identifier for the job (perhaps a UUID or unique string) that we use for tracking status. The other fields are the data needed to perform the task.

2. Enqueue the Job: In the Photon request handler (e.g., an HTTP POST to `/send-email`), you create an `EmailJob` value, assign it a new `id`, serialize it to JSON, and push it onto a Redis list. For example:

```
// Pseudocode for a Photon route handler
fn send_email_handler(req: Request) -> Response {
    // Parse request (e.g., JSON body with email info)
    let email = req.json::<EmailJobData>(); // assume EmailJobData has to,
    subject, body fields

    // Create a job struct with unique ID
    let job_id = generate_unique_id(); // e.g., UUID
    let job = EmailJob {
        id: job_id.clone(),
        to: email.to,
        subject: email.subject,
        body: email.body,
    };

    // Serialize job to JSON
    let payload = serde_json::to_string(&job).expect("Job serialization
failed");

    // Push to Redis queue (e.g., list named "jobs:default")
```

```

let mut redis_conn = get_redis_connection(); // obtain from pool
redis_conn.rpush("jobs:default", payload).expect("Failed to enqueue job");

// (Optional) Record initial status in Redis
redis_conn.hset(format!("job:{}", job_id), "status", "queued").ok();

// Return HTTP 202 Accepted with the job ID
HttpResponse::Accepted()
    .header("Location", format!("/jobs/{}", job_id)) // URL where status
can be checked
    .json(serde_json::json!({ "job_id": job_id }))
}

```

In the above pseudocode:

- We push the JSON to a Redis list called `"jobs:default"` (you could use different queue names for different job types or priorities). This operation is very fast – it's just adding to an in-memory list in Redis – so it typically won't slow down the request noticeably. Photon's coroutine will yield internally while waiting for Redis to acknowledge, so other coroutines can run in the meantime (ensuring the server isn't blocked).
- We also set a Redis hash key `job:<job_id>` with a field `status = "queued"` (this is optional but recommended for tracking). This way, we have a record of the job's status that can be queried later. We use a Redis **hash** (like a map/dictionary) per job to store metadata (status, progress, result, etc.).
- Finally, we return a `202 Accepted` response. A 202 status is appropriate for requests that have been accepted for processing but not yet completed. We include the `job_id` in the response (often in the body) and also a `Location: /jobs/<id>` header pointing to the status endpoint for this job ⁷. The client can use this ID or URL to poll for job status.

3. Do Not Await the Result: Notice that we do not wait for the email to actually be sent. We only waited long enough to enqueue the job in Redis (which is a quick operation). Photon immediately returns the response after enqueueing. This non-blocking fire-and-forget behavior is critical: the user gets an immediate acknowledgement, and the heavy work happens later in a separate process. (If the enqueue itself is slow or if you want extra safety, you could even spawn a Photon coroutine to do the Redis call so the HTTP handler returns instantly. But typically a single Redis push is fine to do inline.)

By following this pattern, **Photon remains responsive** even for endpoints that initiate expensive jobs. We have essentially decoupled job initiation from execution. This approach is confirmed by common practice: for example, a typical flow is to store a job record as PENDING in a database or queue, return 202 Accepted with the job ID, and then let a background worker process it to completion ⁸.

Background Worker Implementation (Rust)

The background workers are separate Rust programs (or threads) that continuously fetch jobs from Redis and execute them. You can run one or many workers depending on throughput needs. Each worker will:

1. **Connect to Redis** and listen on the job queue (e.g., `"jobs:default"` list). This can be done with a blocking pop operation (`BLPOP` or `BRPOP`), which waits until an item is available. Alternatively, you can poll in a loop with a short sleep, but `BLPOP` is more efficient.
2. **Retrieve jobs** in a loop. When a job JSON is retrieved from Redis, parse it (using Serde) into the job struct (e.g., `EmailJob`). At this point, you have the job ID and all data needed.
3. **Update status to "running"** (if using status tracking). For example, set `HSET job:<id> status "running"` in Redis, or update the database record to indicate the job has started. This lets the status endpoint reflect that the job is in progress.
4. **Execute the task.** Perform the actual job logic (send the email, process the data, etc.). This code will depend on the job – it might call an external API, do CPU-intensive work, etc. It should handle errors appropriately.
5. **Update final status.** After completing the job, update the status to "completed", and possibly store the result (if there is a result to report) or any error message if it failed. For example, if using Redis for status: `HSET job:<id> status "completed"` (and maybe `HDEL job:<id> error` if we had an error field). If the job fails (an exception or error is caught), you might set `status = "failed"` and record an `"error"` field with details.
6. **Loop back** to step 2 to process the next job.

A minimal worker loop in Rust (using the `redis` crate for simplicity) might look like:

```
use redis::Commands;
use serde::Deserialize;

// assuming EmailJob struct as defined earlier
let client = redis::Client::open("redis://127.0.0.1/").unwrap();
let mut conn = client.get_connection().unwrap();

loop {
    // Block until a job is available
    let job_json: String = conn.blpop("jobs:default", 0).unwrap().1;
    // .blpop returns (key, value) tuple; 0 means block indefinitely

    // Deserialize the job
    let job: EmailJob = serde_json::from_str(&job_json).expect("Invalid job
JSON");
```

```

let job_key = format!("job:{}", job.id);

// Mark as running
let _ : () = conn.hset(&job_key, "status", "running").unwrap();

// Execute the job task
let result = send_email(&job); // send_email() is the actual function to
perform the email sending

match result {
    Ok(_) => {
        conn.hset(&job_key, "status", "completed").unwrap();
    }
    Err(e) => {
        conn.hset(&job_key, "status", "failed").unwrap();
        conn.hset(&job_key, "error", format!("{}", e)).ok();
    }
}
// loop continues...
}

```

In this snippet, `send_email(&job)` represents the actual work (e.g., using an SMTP client or an API to send the email). The worker marks the job as running before executing, and marks it completed or failed afterward. We use a Redis hash `job:<id>` to keep status (and optionally an `"error"` field if failed).

Parallelism: If you want a single worker process to handle multiple jobs concurrently, you could spawn threads (or coroutines/tasks) within the worker for each job popped. For example, one strategy is to use `BRPOP` with a short timeout in a loop and spawn a thread for each job retrieved. However, a simpler approach is often to run multiple instances of the worker binary (each will do its own BLPOP). Since Redis lists support multiple consumers (each pop gets a different job), you can scale by just running N workers. This is similar to how Sidekiq uses threads – distributing jobs across workers yields parallel processing ⁴.

Job Retries & Reliability: The above simplistic loop assumes jobs are processed at least once. In production, you might want to handle failures more robustly – e.g., if a job fails, you could push it to a retry queue or schedule it to run later. Sidekiq, for instance, will automatically retry failed jobs a few times and keep a "dead" queue for jobs that consistently fail ⁹. You can implement something similar in Rust: for example, have the worker catch errors and use `RPush` to push the job JSON onto a special "retry" list with some delay logic, or use a Redis Sorted Set with timestamps for scheduled retries. Those details depend on your needs. A library like `Apalis` or `queue_workers` crate can handle retries for you ¹⁰ ¹¹, but for a basic implementation you might start without automatic retries and add them later.

Out-of-Process vs In-Process: Notably, Photon's workers are **out-of-process**. This is often preferable for long-running or CPU-intensive tasks, because it isolates the jobs from the web server. (In some frameworks, you could also spawn background tasks in the same process using threads or async tasks – Photon could theoretically spawn an internal coroutine for a job. However, that ties up server resources and risks the job crashing the web process. By keeping workers separate, Photon remains lightweight and focused on handling web requests.)

Tracking Job Status: Redis vs Database

To allow clients (or your frontend/UI) to check the progress of a background job, you need to track job metadata (status, result, etc.) in a place accessible to Photon. There are two common approaches for storing this info: **Redis** or a **database**.

- **Using Redis for Status Tracking:** Since we're already using Redis for the queue, it's natural to use it to track status as well. As shown above, you can maintain a Redis hash per job (e.g. key `job:<id>`). Fields could include `status` (with values like "queued", "running", "failed", "completed"), a `progress` value if the job reports progress, an `error` message if any, or even the final result (if it's small and you want to retrieve it via the API). Storing status in Redis has the advantage of being very fast to read/write and avoids adding another system. Sidekiq itself keeps all job data in Redis for quick access ⁹ (and even provides a web dashboard that reads Redis for job stats ¹²).

Considerations: If you use Redis for status, remember that data in Redis is in-memory (though Redis can persist to disk). **Persistence:** You should enable RDB or AOF persistence in Redis if you want the job statuses to survive Redis or application restarts. Alternatively, if job status being lost on a crash is acceptable (the jobs themselves won't be lost if the queue is persisted, but their last known status might), then Redis alone is fine. Many background job systems treat completed job records as ephemeral (for example, Sidekiq by default doesn't retain completed job entries indefinitely). You might choose to expire (`EXPIRE`) the status keys after some time once jobs are done, to avoid Redis memory growing unbounded.

- **Using a Database for Tracking:** The alternative is to record jobs in a SQL database. This means having a `jobs` table (or multiple tables for different job types) with columns such as `id` (job ID), `type` (job name/type), `status`, `queued_at`, `started_at`, `finished_at`, `result` or `error`, etc. When enqueueing a job, Photon would also insert a row into this table with status "queued". The worker, upon starting the job, would update the row to "running", and finally mark it "completed"/"failed" with any result data or error info. Using a DB has the benefit of **persistence and queryability** – you get a permanent log of jobs. It's useful if you need to display job history, do analytics, or ensure nothing is lost even if Redis is flushed. In the Rust ecosystem, the **Oban** library in Elixir is an example of using a database as a job queue (Oban uses Postgres to store and manage jobs) ¹³, demonstrating that a relational DB can act as both queue and state store for jobs.

Considerations: A downside to tracking in a DB is the extra overhead – every job state change results in a database write, which could be significant if you process many jobs per second. You'll also introduce a coupling between your web app, workers, and the database for job info. If using a DB, ensure that your transactions for enqueue are robust: you want to avoid a scenario where you insert the DB row but fail to push to Redis (or vice versa). One way to handle this is to perform the Redis push **after** the DB insert, and have a periodic reconciling process that checks for any jobs in DB marked queued that might not have been pushed (or use Redis streams with acknowledgements). In practice, this is rarely an issue if the operations happen back-to-back and both succeed most of the time, but it's something to keep in mind for consistency.

On the plus side, a DB makes it easy to query job statuses with complex filters (e.g., "show all failed jobs in the last day") and to integrate with ORMs. Since Photon is using **SeaORM** (an async ORM for Rust) for database access, you can easily define a `Job` entity and use SeaORM to insert and update job records. Each job could have its own ID (which you also embed in the Redis payload) so the worker knows which DB record to update.

Which to choose – Redis or DB? It depends on your requirements:

- If you need **simple, fast, ephemeral** tracking and you're okay with possibly trimming old job data, Redis is often sufficient. It avoids adding more load to your primary database. This is ideal for high-volume background tasks where you mostly care about recent statuses and can discard old data or move it elsewhere if needed.
- If you need **durability and rich querying** (audit logs, complex admin queries, long-term history), a database is more appropriate. Also, if your web framework already heavily uses a DB and you want everything in one place, this might be convenient.

Some frameworks even allow switching between these backends. For instance, Loco (a Rust web framework) supports both a Redis-backed queue and a Postgres-backed queue for jobs ¹⁴. The job logic in your app remains the same; only the storage changes. This demonstrates that using Redis vs using a DB are both valid patterns, and you can choose based on needs. It's also possible to use **both**: for example, use Redis to queue and track current jobs, but periodically archive completed jobs to a DB for long-term storage (this could be a separate cleanup worker that moves data from Redis to DB).

Example: Minimal Job and Worker Code

To tie it together, let's sketch a minimal implementation for a **queue + worker** system in Rust, using JSON over Redis:

Job Definition

We already defined our `EmailJob` struct and made it `Serialize`/`Deserialize`. In a real setup, you might have multiple job types. You could handle that by having a single enum with variants for each job type, or by having separate Redis queues per job type (e.g., `"queue:emails"`, `"queue:reports"`, etc.) and separate worker binaries for each type. For simplicity, we'll assume one type here.

If you had multiple types in one queue, you might include a `job_type` field in the JSON to know how to process it. Sidekiq, for example, includes a `"class"` name in the job JSON to indicate which worker logic to run ². In Rust, without a dynamic class loader, you'd likely implement a dispatch in the worker: e.g., examine a `task` field and call the appropriate function for that task.

Enqueuing Code (Photon Side)

We showed a pseudocode example in the **Enqueuing Jobs** section for adding a job to Redis. Here's a more concrete snippet, assuming you have a `redis::Connection` (synchronous):

```
use redis::Commands;
use uuid::Uuid;

fn enqueue_email_job(redis_conn: &mut redis::Connection, to: &str, subject:
&str, body: &str) -> String {
```

```

let job_id = Uuid::new_v4().to_string();
let job = EmailJob {
    id: job_id.clone(),
    to: to.to_string(),
    subject: subject.to_string(),
    body: body.to_string(),
};
let payload = serde_json::to_string(&job).unwrap();
// Push to Redis list
redis_conn.rpush("jobs:default", payload).unwrap();
// Set initial status
redis_conn.hset(format!("job:{}", job_id), "status", "queued").unwrap();
job_id
}

```

This function creates a new `EmailJob`, serializes it, pushes it onto `"jobs:default"`, and records the status. It returns the generated `job_id`. In a Photon HTTP handler, you would call this function and then return a response with the `job_id` (as shown earlier). If this operation might be slow or blocking in Photon's coroutine context, you could wrap it in a `may::coroutine::spawn` to run concurrently. But generally, as long as you have a Redis connection ready, `rpush` / `hset` are quick operations.

Note: It's good practice to use a connection **pool** for Redis (and database) connections, especially in a web server that handles many requests. Photon can initialize a pool of Redis connections (there are libraries like `r2d2_redis` or you might use **Lifeguard** if it's an object pool manager, to pool Redis connections as objects). Then each request handler can grab a connection from the pool to enqueue the job, avoiding the overhead of reconnecting to Redis every time. Photon's coroutine model can work with such pooling similarly to threaded models – just ensure proper synchronization when borrowing connections. Lifeguard, being a generic object pool, could manage these connections or any other resources your coroutines need to access in a controlled way.

Worker Code

We outlined a basic worker loop above. Here's how a standalone worker program might look using an async runtime (Tokio) and the asynchronous Redis API, which could integrate well if also using SeaORM in the worker:

```

use redis::AsyncCommands;
use tokio::time::sleep;
use std::time::Duration;
use serde::Deserialize;
use sea_orm::{DatabaseConnection, EntityTrait}; // if using SeaORM for DB
updates

#[derive(Deserialize)]
struct EmailJob { /* fields as before */ }

```



```

#[tokio::main]
async fn main() -> redis::RedisResult<()> {
    let client = redis::Client::open("redis://127.0.0.1/");
    let mut conn = client.get_async_connection().await?;

    // If using DB:
    let db: DatabaseConnection = /* initialize SeaORM DB connection */;

    loop {
        // Use BLPOP with a timeout to get a job or None if no job in timeout
        window
        let job_json: Option<String> = conn.blpop("jobs:default", 5).await?;
        if job_json.is_none() {
            // No job found in 5 seconds, pause or continue loop
            continue;
        }
        let job_json = job_json.unwrap();
        let job: EmailJob = serde_json::from_str(&job_json)?;

        // Mark as running
        let job_key = format!("job:{}", job.id);
        let _: () = conn.hset(&job_key, "status", "running").await?;
        // (If using DB, update DB record to status "running")
        // JobEntity::update(...).exec(&db).await?;

        // Execute the job task
        let outcome = send_email_async(&job).await;

        match outcome {
            Ok(_) => {
                conn.hset(&job_key, "status", "completed").await?;
                // Update DB status to completed if using DB
            }
            Err(e) => {
                conn.hset(&job_key, "status", "failed").await?;
                conn.hset(&job_key, "error", format!("{}", e)).await?;
                // Update DB status to failed, record error if using DB
            }
        }
    }
}

```

This example uses Tokio for async. It fetches a job with `blpop` with a timeout (to allow breaking the loop gracefully if needed – you could also use `BLPOP key 0` to block indefinitely). It updates Redis status and calls an async `send_email_async` function to do the work (which could, for example, use an async HTTP client to call an email service). After the job, it sets the final status in Redis (and likewise in DB if applicable).

You would compile this as a separate binary, e.g., `photon_worker`. You can run multiple instances of `photon_worker` to scale up consumers. In a production deployment, you might run, say, 4 worker processes alongside your Photon web server processes.

Sharing Code: It's wise to put the job struct definitions and any shared logic in a common library crate that both Photon and the worker can depend on. For instance, both need the `EmailJob` struct (Photon to serialize and enqueue it, the worker to deserialize and execute it). By sharing a crate, you ensure the struct and serialization format remain consistent. Similarly, if using SeaORM, you can define your `Job` Entity in one place and use it in both binaries.

Job Status Endpoint Design

Photon should expose an endpoint (or multiple endpoints) to query the status (and possibly result) of a job by its ID. A typical design is a **GET** request to something like `/jobs/<job_id>` that returns a JSON with the job's information.

For example, the endpoint could be `GET /jobs/{id}` and the response might look like:

```
{
  "id": "abcd-1234-efgh-5678",
  "status": "completed",
  "result": "...",
  "error": null,
  "progress": null
}
```

The fields `result`, `error`, `progress` are optional depending on your use case: - `result` could hold the outcome of the job if there's something to return (for instance, if the job generates a report file, the result might be a URL to download it, or some identifier). - `error` would contain an error message if the job failed. - `progress` can be a percentage or description of progress if the job updates progress as it goes (e.g., "42%" or "step 3/5 completed"). To support this, your worker can periodically update a `progress` field in the Redis hash or DB as it works. For long jobs, this can be very useful for UX.

Implementing the status endpoint: In the handler for `GET /jobs/{id}`, you will do roughly:

- Parse the `{id}` from the URL.
- Fetch the job status info from the tracking store. If using Redis, you would retrieve the hash `job:<id>` (e.g., using `HGETALL` or specific `HGET` fields for status, result, etc.). If using a DB, you would query the `jobs` table for that ID (using SeaORM's finder or raw SQL).
- If the job exists, format and return the data as JSON. If not, return 404 Not Found.

Example in pseudocode (for Redis):

```

fn get_job_status(redis_conn: &mut redis::Connection, job_id: &str) ->
HttpResponse {
    let key = format!("job:{}", job_id);
    match redis_conn.hget(&key, "status") {
        Ok(Some(status)) => {
            let status: String = status;
            // you could get more fields like error, result similarly
            let error: Option<String> = redis_conn.hget(&key,
"error").unwrap_or(None);
            let progress: Option<String> = redis_conn.hget(&key,
"progress").unwrap_or(None);
            HttpResponse::Ok().json(serde_json::json!({
                "id": job_id,
                "status": status,
                "error": error,
                "progress": progress
            })))
        }
        Ok(None) => HttpResponse::NotFound().body("Job ID not found"),
        Err(err) => {
            eprintln!("Redis error: {}", err);
            HttpResponse::InternalServerError().body("Error checking job
status")
        }
    }
}

```

In a real implementation, you'd likely structure this with Photon's routing system. But essentially, it's a lookup followed by a conditional response.

If using a database, the logic would be similar: query by ID using SeaORM, and return the fields. For example, using SeaORM's ActiveModel or Entity:

```

if let Some(job) = Job::find_by_id(job_id).one(&db).await? {
    return HttpResponse::Ok().json(JobStatus {
        id: job.id,
        status: job.status,
        error: job.error,
        progress: job.progress,
        // ... etc
    });
} else {
    return HttpResponse::NotFound().body("No such job");
}

```

(Here `JobStatus` could be a struct to serialize as JSON for output.)

Security & Access Control: If these jobs are triggered by user actions, you may want to ensure that only the user who initiated a job can see its status (or an admin). This might require associating a user ID with the job and checking it. That goes beyond the technical implementation, but keep it in mind for real applications.

Polling vs Webhooks: The simplest way to check status is polling this endpoint periodically from the client. Alternatively, one could design a callback or webhook system where the server calls back or pushes an update when the job is done. That's more complex (it requires either client long-polling, WebSockets, or server-to-server callbacks). In many cases, a simple polling of `/jobs/{id}` every few seconds is enough. The `202 Accepted` response from the initial request can hint the client to poll at the status URL ⁷.

Often, the initial POST request that enqueued the job returns `202 Accepted` and includes the URL of the status endpoint (as we did with the `Location` header). For example, the HTTP response might be:

```
HTTP/1.1 202 Accepted
Location: /jobs/abcd-1234-efgh-5678
Content-Type: application/json

{ "job_id": "abcd-1234-efgh-5678" }
```

This informs the client where to check next. The status endpoint itself (when the job is still running) should return `200 OK` with the current status payload. (In other words, once the job is created, **the status endpoint is a resource that always exists** – it's just that its `status` field will change from "queued" to "running" to "completed". You wouldn't return 202 from the GET; you reserve 202 for the POST that initiates the async process.)

Integration with Photon, Lifeguard, and SeaORM

Finally, a few notes on integrating this system into Photon's architecture:

- **Coroutine (may) Compatibility:** Photon's use of the `may` coroutine library means it's not built on Tokio's async runtime. However, you can still use Redis and even async libraries like SeaORM; you just have to manage how they are driven. If Photon's runtime can perform socket I/O without blocking (e.g., using `may::net` for non-blocking sockets), the `redis` crate's connection might block unless you use non-blocking mode. One approach is to use a threadpool for blocking I/O. For example, Photon could use something like `may::sync::defer` or simply spawn a dedicated thread (from a pool) to handle each Redis or DB operation, so that the coroutine yields while that thread does the work. If **Lifeguard** is being used, it might serve as an object pool for such threads or connections, ensuring that expensive resources are reused and that Photon's coroutines don't stall each other.
- **Database Access (SeaORM):** SeaORM is an async ORM (built on SQLx). In a purely synchronous environment, calling async DB methods would require a runtime. You can resolve this by running a small Tokio runtime in a separate thread for DB operations or by using SeaORM's connection in a

blocking way. Another option is to use SeaORM's **runtime compatibility** feature – SQLx (which SeaORM uses under the hood) can run on Tokio or on any runtime that implements the async I/O traits. If Photon's `may` has an async I/O compatibility layer, you might integrate it. If not, the simplest route is to handle DB operations outside of Photon's coroutine threads (again, perhaps via a threadpool). For instance, Photon could enqueue a DB insert or update to a channel that a separate thread (running Tokio) listens on. That thread executes the DB operation and perhaps returns the result if needed. This is similar to how one might integrate blocking code in an async server (like using `spawn_blocking` in Tokio).

- **Lifeguard:** As an object pool manager, Lifeguard can manage pools for both Redis connections and database connections. Photon can initialize a pool of Redis connections at startup (minimize overhead of connecting to Redis for each job enqueue). Similarly, a pool of database connections (or SeaORM DatabaseConnection, which itself often manages a pool internally) can be prepared. Whenever a coroutine handler needs a connection, it can borrow from the pool quickly (this is typically a non-blocking constant-time operation), do the work, and return it. Lifeguard ensures that if all connections are busy, requests wait rather than overwhelming the DB or Redis with new connections. Essentially, Lifeguard + SeaORM could play a role similar to `r2d2 + Diesel` in a synchronous environment.
- **Job Metadata Schema:** If using SeaORM, define an Entity for jobs. For example:

```
#[derive(Debug, Clone, PartialEq, DeriveEntityModel)]
#[sea_orm(table_name = "jobs")]
pub struct Model {
    #[sea_orm(primary_key)]
    pub id: String,
    pub job_type: String,
    pub status: String,
    pub error: Option<String>,
    pub result: Option<String>,
    pub queued_at: DateTime,
    pub started_at: Option<DateTime>,
    pub finished_at: Option<DateTime>,
}
```

(Add fields as needed, e.g., progress, payload, etc.)

You can then use `Job::insert` when enqueueing (with status "queued"), and `Job::update` in the worker to set status "running" and "finished". SeaORM will handle the SQL. Just ensure both Photon and worker have access to run these queries (likely via a shared codebase or crate with the Entity defined). Photon's DB pool (managed by Lifeguard/SeaORM) would be used for inserts, and the worker can have its own DB connection (pool) for updates.

- **Error Handling and Monitoring:** Integrating with Photon also means considering how you'll monitor this background job system. You might want Photon to expose an admin endpoint (or UI) listing recent jobs, which would query either Redis or the DB. If using Redis only, you might build an

endpoint that scans keys by pattern (like `job:*`) – but that can be inefficient if there are many jobs. If using a DB, you can query with limits and filters easily. Logging is also important: workers should log failures so you can debug issues. If a job keeps failing, you might have Photon provide a way to see that error (via the status endpoint or a separate error queue).

- **Cancellation:** Though not asked, consider how you might allow canceling a job. If a job is queued (not yet running), you could remove it from the Redis list (though Redis lists don't have a direct remove by value, you could maybe use `LREM` if you store the exact JSON – or maintain a separate index). If a job is running, cancellation is harder – you would need the worker to check a flag (for example, if you set `status = "cancelled"`, a cooperative worker could see that and stop early). Cancellation is a complex topic, but designing the status tracking with a possible "cancelled" state might be worthwhile if your use case calls for it.

In summary, **Photon's coroutine model** works very well with this out-of-process job approach: the web server stays lightweight, doing just an enqueue. The heavy work is done by separate Rust worker processes that can be scaled as needed ⁴. Using **Redis** as the communication medium provides a fast, persistent queue (jobs won't be lost on restart because Redis can persist data to disk – a necessary feature for reliability ¹⁵). For **job status tracking**, you have a design choice between using Redis (quick and simple) or a **database** (robust and queryable with SeaORM) – or even both in combination. The implementation can be achieved with pure Rust easily (with crates like `redis`, `serde_json`, etc.), or you can leverage higher-level libraries (like Apalis, Sidekiq-rs, or `queue_workers`) that provide some of this functionality out of the box.

By following this architecture, Photon can handle background jobs much like a Rails app with Sidekiq: clients fire off requests that enqueue jobs, get immediate responses, and then poll a status endpoint to see when the work is done. This decoupling makes your web application scalable and your background processing reliable and trackable. **In code, it means writing a bit of enqueue logic in Photon and a simple Rust worker program – as illustrated above – to perform the jobs.** With coroutines and tools like Lifeguard and SeaORM, integrating these pieces should be straightforward and performant.

Sources:

- Sidekiq background jobs and Redis (for conceptual model) ¹ ²
- Stack Overflow – patterns for long-running job APIs (202 Accepted and status endpoint) ⁷
- Rust background job frameworks (Apalis, etc.) and usage of Redis/DB ⁶ ¹³
- Considerations on task persistence and async processing in Rust ¹⁶

¹ ² ³ ⁴ ⁹ ¹² Sidekiq and Redis: A Powerful Partnership for Background Jobs in Ruby | by Hardik | Medium

<https://medium.com/@radadiyahardik355/sidekiq-and-redis-a-powerful-partnership-for-background-jobs-in-ruby-519ae135efa9>

⁵ ⁶ Background job processing with rust using actix and redis

<https://mureithi.me/blog/background-job-processing-with-rust-actix-redis>

⁷ HTTP Status Code for "Still Processing"

<https://softwareengineering.stackexchange.com/questions/316208/http-status-code-for-still-processing>

8 java - How to monitor REST Endpoint for long running jobs - Stack Overflow

<https://stackoverflow.com/questions/56923600/how-to-monitor-rest-endpoint-for-long-running-jobs>

10 11 queue_workers — async Rust library // Lib.rs

https://lib.rs/crates/queue_workers

13 15 16 Fang, async background processing for Rust - Kraken of Thought

<https://badykov.com/rust/async-processing/>

14 Workers - Loco.rs

<https://loco.rs/docs/processing/workers/>